

Part (a):

After pre-processing the data, checking and handling all the missing values, taking care of categorical and ordinal features and null values I ended up with 46 features that according to me should be important for the model prediction.

So, after pre-processing the data from 111 features I have shrunk my dataset down to 46 features only.

Part (b):

Accuracy, Precision and Recall for each of the model

Model 1:

Model 1 (n_estimators=50,max_features=5,max_depth=5)

```
In [53]: start_time = time.time()

model1 = GradientBoostingClassifier(n_estimators=50,max_features=5,max_depth=5)
model1.fit(X_train, y_train)
pred = model1.predict(X_test)
print('Accuracy Score: ',accuracy_score(y_test, pred)*100)

end_time=time.time()
print("Time Taken: {} secs".format(round(end_time-start_time,3)))
print('\n')
# Checking classification_report
print('Classification Report')
print('\n')
print(classification_report(y_test, pred))
```

Accuracy Score: 92.96041549691185
Time Taken: 1.566 secs

Classification Report

	precision	recall	f1-score	support
+1	0.93	0.99	0.96	12122
-1	0.93	0.57	0.71	2126
accuracy			0.93	14248
macro avg	0.93	0.78	0.83	14248
weighted avg	0.93	0.93	0.92	14248

Accuracy : 92.96 %

Precision and Recall :

	precision	recall
+1	0.93	0.99
-1	0.93	0.57

Model 2:

Model 2 (n_estimators=100,max_features=5,max_depth=5)

```
In [54]: start_time = time.time()

model2 = GradientBoostingClassifier(n_estimators=100,max_features=5,max_depth=5)
model2.fit(X_train, y_train)
pred = model2.predict(X_test)
print('Accuracy Score: ',accuracy_score(y_test, pred)*100)

end_time=time.time()
print("Time Taken: {} secs".format(round(end_time-start_time,3)))
print('\n')
# Checking classification_report
print('Classification Report')
print('\n')
print(classification_report(y_test, pred))
```

Accuracy Score: 95.12914093206064
Time Taken: 2.937 secs

Classification Report

	precision	recall	f1-score	support
+1	0.96	0.99	0.97	12122
-1	0.92	0.74	0.82	2126
accuracy			0.95	14248
macro avg	0.94	0.86	0.90	14248
weighted avg	0.95	0.95	0.95	14248

Accuracy : 95.12 %

Precision and Recall :

	precision	recall
+1	0.96	0.99
-1	0.92	0.74

Model 3:

Model 3 (n_estimators=150,max_features=5,max_depth=5)

```
In [55]: start_time = time.time()

model3 = GradientBoostingClassifier(n_estimators=150,max_features=5,max_depth=5)
model3.fit(X_train, y_train)
pred = model3.predict(X_test)
print('Accuracy Score: ',accuracy_score(y_test, pred)*100)

end_time=time.time()
print("Time Taken: {} secs".format(round(end_time-start_time,3)))
print('\n')
# Checking classification_report
print('Classification Report')
print('\n')
print(classification_report(y_test, pred))
```

Accuracy Score: 95.66956765861875
Time Taken: 4.492 secs

Classification Report

	precision	recall	f1-score	support
+1	0.96	0.99	0.97	12122
-1	0.91	0.79	0.85	2126
accuracy			0.96	14248
macro avg	0.93	0.89	0.91	14248
weighted avg	0.96	0.96	0.96	14248

Accuracy : 95.66 %

Precision and Recall :

	precision	recall
+1	0.96	0.99
-1	0.91	0.79

Final Model:

Final Model (n_estimators=500,max_features=8,max_depth=10)

n_estimators = 500, max_features = 8, max_depth = 10 leads to the best accuracy

```
In [56]: start_time = time.time()

final_model = GradientBoostingClassifier(n_estimators=500,max_features=8,max_depth=10)
final_model.fit(X_train, y_train)
pred = final_model.predict(X_test)
print('Accuracy Score: ',accuracy_score(y_test, pred)*100)

end_time=time.time()
print("Time Taken: {} secs".format(round(end_time-start_time,3)))
```

Accuracy Score: 96.80656934306569
Time Taken: 54.273 secs

Best test accuracy achieved is \approx 97%

```
In [57]: # Checking classification_report
print('Classification Report')
print('\n')
print(classification_report(y_test, pred))
```

Classification Report

	precision	recall	f1-score	support
+1	0.97	0.99	0.98	12122
-1	0.93	0.85	0.89	2126
accuracy			0.97	14248
macro avg	0.95	0.92	0.93	14248


Accuracy : 96.80 %

Precision and Recall :


	precision	recall
+1	0.97	0.99
-1	0.93	0.85

So, Best Accuracy achieved is : 96.80%

Checking for the Important features having atleast 1% corr with our class label as per our model



Effect of increasing the number of trees on the test error:



As we increase the number of trees the test error starts to decrease and it will come to a point where increasing the number of trees any further will not affect the test error (like from 400-500 trees in the diagram above).

Comparing the best performance (accuracy, precision, recall) of our model against a simple decision tree built using information gain:

Using GradientBoosting Classifier(Best One):

Final Model (n_estimators=500,max_features=8,max_depth=10)

n_estimators = 500, max_features = 8, max_depth = 10 leads to the best accuracy

```
In [56]: start_time = time.time()

final_model = GradientBoostingClassifier(n_estimators=500,max_features=8,max_depth=10)
final_model.fit(X_train, y_train)
pred = final_model.predict(X_test)
print('Accuracy Score: ',accuracy_score(y_test, pred)*100)

end_time=time.time()
print("Time Taken: {} secs".format(round(end_time-start_time,3)))
```

Accuracy Score: 96.80656934306569
Time Taken: 54.273 secs

Best test accuracy achieved is ≈ 97%

```
In [57]: # Checking classification_report
print('Classification Report')
print('\n')
print(classification_report(y_test, pred))
```

Classification Report

	precision	recall	f1-score	support
+1	0.97	0.99	0.98	12122
-1	0.93	0.85	0.89	2126
accuracy			0.97	14248
macro avg	0.95	0.92	0.93	14248

Accuracy: ~ 97 %

Precision and Recall:

	precision	recall
+1	0.97	0.99
-1	0.93	0.85

Simple Decision Tree built using Information Gain:

Comparing GradientBoosting Performance against a simple decision tree

```
In [63]: #simple decision tree built using information gain
dt_model = DecisionTreeClassifier(criterion = 'entropy')
dt_model.fit(X_train,y_train)
pred = dt_model.predict(X_test)
accuracy_score(pred,y_test)
```

```
Out[63]: 0.9390791690061763
```

Test accuracy achieved using decision tree is $\approx 94\%$

```
In [64]: # Checking classification_report
print('Classification Report')
print('\n')
print(classification_report(y_test, pred))
```

Classification Report

	precision	recall	f1-score	support
+1	0.96	0.96	0.96	12122
-1	0.80	0.79	0.79	2126
accuracy			0.94	14248
macro avg	0.88	0.88	0.88	14248
weighted avg	0.94	0.94	0.94	14248

Accuracy: $\sim 94\%$

Precision and Recall:

	precision	recall
+1	0.96	0.96
-1	0.80	0.79